



Memory Management

Dr Tran Duc Minh and Dr Hung Tran

DATCOM Lab
Faculty of Data Science and Artificial Intelligence
College of Technology, National Economics University
Email: minhdt@neu.edu.vn, hung.tran@neu.edu.vn

April 21, 2025

Outline

- 1 Structure of the Page Table
- 2 Swapping




Hierarchical Paging

- Most modern computer systems support a large logical address space (2^{32} to 2^{64}).
- In such an environment, the page table itself becomes excessively large.
- For example: Consider a system with a 32-bit logical address space.
 - ▶ If the page size is 4 KB (2^{12}), then a page table may consist of over 1 million entries ($2^{20} = 2^{32}/2^{12}$).
 - ▶ Assuming that each entry consists of 4 bytes, each process may need up to 4 MB of physical address space for the page table alone.
- We would not want to allocate the page table contiguously in main memory.
- One simple solution to this problem is to divide the page table into smaller pieces.
- We can accomplish this division in several ways.




Two-Level Paging Algorithm



Two-Level Paging Algorithm

- Consider a system with a 32-bit logical address space and a page size of 4 KB.
- A logical address is divided into a page number (20 bits) and a page offset (12 bits).
- Since we page the page table, the page number is further divided into a 10-bit page number and a 10-bit page offset. Thus, a logical address is structured as follows:




where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table.

- Since address translation works from the outer page table inward, this scheme is known as a **forward-mapped** page table.



Two-Level Paging Algorithm



Two-Level Paging in 64-bit Systems


- For a system with a 64-bit logical address space, a two-level paging scheme is no longer appropriate.
- With a page size of 4 KB (2^{12}), the page table consists of up to 2^{52} entries.
- In a two-level paging scheme, the inner page tables can contain 2^{10} 4-byte entries.
- The outer page table would consist of 2^{42} entries, or 2^{44} bytes.

outer page	inner page	offset
p_1	p_2	d
42	10	12



Three-Level Paging in 64-bit Systems

- For example: We can page the outer page table, giving a three-level paging scheme.
 - Suppose the outer page table is made of standard-size pages (2^{10} entries, or 2^{12} bytes).
 - In this case, a 64-bit address space still results in an outer page table of 2^{34} bytes (16 GB) in size.
 - And possibly 4 memory access to get to one physical memory location




Hashed Page Table for Large Address Spaces

- Address spaces larger than 32 bits is to use a **hashed page table** with the hash value is the **virtual page number**.
- Each entry in the hash table contains a linked list of elements that hash to the same location (to handle collisions).
- Each element consists of three fields:
 - ① The virtual page number,
 - ② The value of the mapped page frame,
 - ③ A pointer to the next element in the linked list.
- The algorithm works as follows:
 - ▶ The virtual page number in the virtual address is hashed into the hash table.
 - ▶ The virtual page number is compared with field 1 in the first element of the linked list.
 - ▶ If there is a match, the corresponding page frame (field 2) is used to form the desired physical address.
 - ▶ If no match, subsequent entries in the linked list are searched for a matching virtual page number.



Hashed Page Table for Large Address Spaces



Clustered Page Tables for 64-bit Systems

- A variation of the hashed page table scheme that is useful for 64-bit address spaces.
- This variation uses **clustered page tables**, which are similar to hashed page tables except that each entry in the hash table refers to several pages (such as 16) rather than a single page.
- A single page-table entry can store the mappings for multiple physical-page frames.
- Clustered page tables are particularly useful for **sparse** address spaces, where memory references are noncontiguous and scattered throughout the address space.




Inverted Page Tables

- **Traditional Page Tables:** One entry per virtual page per process, leading to high memory usage due to many unused entries.
- **Inverted Page Table:**
 - ▶ One entry per physical frame, not per virtual page.
 - ▶ Each entry stores the virtual page number and process ID (PID).
 - ▶ Size depends on physical memory, not virtual address space.
- **Advantages:** Reduces memory consumption for page tables.
- **Operation:** OS searches the table to map virtual addresses to physical frames.
- **Usage:** Common in systems like 64-bit UltraSPARC and PowerPC, often with an address-space identifier.
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs



Inverted Page Tables




Standard swapping

- **Purpose:** Allows a process (or part) to be temporarily moved from memory to a backing store and back for execution.
- **Benefit:** Enables total physical address space of processes to exceed physical memory, increasing multiprogramming.
- **Standard Swapping:**
 - ▶ Moves entire processes between main memory and a **Backing store**.
 - ▶ Backing store must be large, provide direct access, and store process data structures (including per-thread data for multithreaded processes).
- **OS Role:** Maintains metadata for swapped-out processes to restore them when swapped back in.
- **Advantages:**
 - ▶ Allows memory oversubscription to accommodate more processes than physical memory can hold.
 - ▶ Frees memory from idle processes for active ones.
- **Process:** Idle processes are swapped out; if they become active, they are swapped back in.



Standard swapping




Swapping with Paging

- Modern systems (Linux, Windows) swap individual pages, not entire processes.
- More efficient: only a small number of pages are swapped, still allows memory oversubscription.
- **Terminology:**
 - ▶ "Swapping" refers to standard swapping.
 - ▶ "Paging" refers to swapping with paging.
- **Operations:**
 - ▶ *Page out*: Moves a page to the backing store.
 - ▶ *Page in*: Retrieves a page from the backing store.



Standard swapping



Swapping on Mobile Systems

- **Mobile Systems:** Avoid swapping due to flash memory constraints:
 - ▶ Limited storage space compared to hard disks.
 - ▶ Poor throughput between main memory and flash memory.
- **iOS Strategy:**
 - ▶ Requests apps to voluntarily free memory when low.
 - ▶ Read-only data (e.g., code) can be removed and reloaded from flash memory.
 - ▶ Modified data (e.g., stack) are never removed.
 - ▶ Apps failing to free enough memory may be terminated.
- **Android Strategy:**
 - ▶ Similar to iOS: may terminate processes if memory is low.
 - ▶ Saves application state to flash memory before termination for quick restarts.
- **Developer Responsibility:** Mobile developers must carefully manage memory to avoid excessive usage or leaks.



Questions

